

EIS: The EUV Imaging Spectrometer on Hinode

EISPAC USER'S GUIDE

EISPAC: EIS Python Analysis Code

Harry P. Warren, Micah J. Weberg, Nicholas A. Crump, and the EIS Team
Space Science Division, Naval Research Laboratory, Washington, DC 20375

Version 0.7 | January 31, 2021

This software development was sponsored by NASA's *Hinode* Program

DISTRIBUTION STATEMENT A. Approved for public release: distribution unlimited

Contents

<i>Introduction</i>	3
<i>Downloading and Reading the Data</i>	6
<i>Fitting the Data</i>	11
<i>Level-1 HDF5 File Details</i>	16
<i>Acknowledgments</i>	19
<i>Appendix A: mpfit documentation</i>	20
<i>Appendix B: Useful EISPAC doc strings</i>	22
<i>Bibliography</i>	25

Introduction

THE EUV IMAGING SPECTROMETER — EIS — was designed to study the solar atmosphere and answer fundamental questions on the heating of the solar corona, the origin of the solar wind, and the release of energy in solar flares¹. EIS observes two wavelength ranges in the extreme ultraviolet, 171—212 Å and 245—291 Å with a spectral resolution of about 22 mÅ and a plate scale of 1'' per pixel. Solar images can be made by stepping the slit over a region of the Sun and taking an exposure at each position. A detailed description of EIS is given in the instrument paper².

This document describes the basic elements of EIS data analysis using new HDF5 level-1 files and the EIS Python Analysis Code (EISPAC) package. At the beginning of the Hinode mission the strategy was to release unprocessed level-0 FITS files and software routines written in IDL for processing these files into a format that could be used for data analysis. Additionally, all of the routines for computing ancillary information, such as the offsets of the detectors or the magnitude of the instrumental broadening, were all written in IDL. Unfortunately, IDL is an expensive, proprietary language, little used outside of solar physics. Python, in contrast, is a free, open source language that has grown dramatically in popularity since the launch of Hinode, making it an obvious choice for future software development.

EIS Level-1 HDF5 Files

To accelerate the transition to Python we have created a new level-1 product that contains both the processed level-1 data and the ancillary information needed for data analysis. The alternative approach, to port all of the existing IDL software to Python, would be time consuming and create confusion about which routines are being actively supported during the transition. Distributing level-1 files removes this problem, but does make the user dependent on the team for reformatting all of the files as bugs are discovered. Since the mission has been going on for some time now, the number of bugs is likely to be small.

There are several other design decisions that merit some explanation

- The data and header information are stored in separate files. Since

¹ EIS is part of the Hinode mission and was sponsored by the Japan Aerospace Exploration Agency (JAXA), the United Kingdom Space Agency (UKSA), and National Aeronautics and Space Administration (NASA) with contributions from ESA and Norway. Hinode was launched on September 22, 2006 at 21:36 UTC from the Uchinoura Space Center in Japan and continues to operate.

² Culhane, J. L., Harra, L. K., James, A. M., et al. 2007, *Sol. Phys.*, 243, 19

the data is large and unlikely to change, the time-consuming download of these files should only need to be done once. The header file is very small and can be updated easily.

- HDF5 is used to store the data. This is a very widely used, high-performance file format that is well supported by both IDL and Python. The most attractive feature for this application is that data is stored in a self-documenting, directory-like tree structure instead of binary table extensions.
- The data is processed from raw “data numbers” to “photon events” or “counts”. The default behavior of `eis_prep` is to convert to calibrated units. With the HDF5 files conversion to absolute units is done using a calibration curve in the header file, and several different calibration curves can be considered.

The processed level-1 HDF5 files can be downloaded either directly from the NRL Hinode/EIS website at <https://eis.nrl.navy.mil/> or using the tools included in EISPAC. Chapter 4 describes the processing of the files in more detail.

EIS Python Analysis Code (EISPAC)

EISPAC provides Python classes and functions that can read the new HDF5 files, perform all of the necessary calibration and pointing adjustments, and create user-friendly Python objects that can be manipulated as needed. Also included are functions for fitting the intensity profiles with multi- Gaussian functions using template files and a Python port of the venerable MPFIT library (Markwardt, 2009). For convenience, command line and GUI tools are provided to help users quickly browse and download data, copy template files, and fit multiple files at once using parallel processing.

Requirements

EISPAC depends on a number of Python packages that are commonly used in scientific and solar research. Normally, the installation process should automatically check and install missing dependencies, assuming your environment is configured appropriately. If it does not, you may wish to try installing the required packages individually first.

- `python` \geq 3.7
- `numpy` \geq 1.18.1
- `scipy` \geq 1.4.1
- `matplotlib` \geq 3.1

- `h5py` \geq 2.9
- `astropy` \geq 3.1
- `sunpy` \geq 1.0.3
- `ndcube` \geq 1.2.1
- `pyqt` \geq 5.9

Additionally, some of the command line tools in EISPAC depend on two non-Python software packages - `wget` and `cURL`. Both of these packages should come preinstalled on most modern operating systems. If your system does not, please refer to the respective project websites and/or contact your system administrator.

Installation

This initial release of EISPAC is not yet available on the usual Python web repositories (PyPi or Conda). As such, the installation process is a little bit more involved than other packages.

1. Download the entire "eispac_develop" repository and extract it to a convenient directory on your computer (it does not matter where).
2. Open a terminal and navigate to the directory chosen above
3. Run the install script using your preferred package manager,
 - (a) **PIP**: type `python -m pip install .`
 - (b) **conda**: DETAILS FORTHCOMING

If you later wish to update EISPAC you will need to repeat steps 1 & 2 above and then issue the command `python -m pip install --upgrade .`
You should be all ready to go now!

Downloading and Reading the Data

There are two main components to the EISPAC software: (1) a set of command line scripts and (2) the Python package itself. In this chapter we will give a very brief overview of how to read and explore the EIS data contained in the level-1 HDF5 files. The next chapter covers how to fit the data with Gaussian functions.

Using the command line scripts

The command line scripts should be automatically installed and registered with the OS as part of installing EISPAC. These scripts are designed to help users quickly browse, download, and fit Gaussian functions to the data. To use a script, simply enter its name in the command line from any directory in which you have read and write privileges.

There are currently four command line scripts available,

eis_catalog - GUI tool from searching the as-run EIS data catalog and downloading the HDF5 files your computer. Can also generate a text list of files to download.

eis_browse_templates - GUI tool for browsing the fit templates corresponding to each spectral window in a given observation set and copying the template files from EISPAC to your current working directory (fit templates are explained more in the next chapter)

eis_download_files - Command line tool for downloading a the level-1 HDF5 files assoiated with one or more level-0 EIS fits files. Can also download an entire list of files using the text output of `eis_catalog`. Example usage,

```
>>> eis_download_files eis_l0_20190404_131513.fits
```

eis_fit_files - Command line tool for fitting all of the HDF5 files in a given directory with each fit template found in another directory. Example usage,

```
>>> eis_fit_files ./eis_study/ ./eis_study/templates/
```

Note well: some scripts will default to saving files to your current working directory, therefore we recommend running the scripts from the directory in which you intend to do most of your analysis.

Reading and Exploring data with EISPAC

Once installed, EISPAC can be imported into any Python script or interactive session with a simple `import eispac` statement. Assuming you have already downloaded some data, the following code snippet below illustrates how to read the level-1 data from the spectral window containing the Fe XII 195.12 Å line (window 7, in our example file). At the end of the chapter we will show how to examine a data header file to determine what wavelengths are available in a given observation.

```
>>> import eispac
>>> data_filename = 'eis_20190404_131513.data.h5'
>>> data_cube = eispac.read_cube(data_filename, 195.12)
```

The `read_cube()` function will read and apply all of the calibration and pointing corrections necessary for scientific analysis. The function takes three arguments:

filename (str or pathlib path) - Name or path of either the data or header HDF5 file for a single EIS observation

window (int or float, optional) - Requested spectral window number (if ≤ 24) or the value of any wavelength within the requested window (in units of [Å]). Default is "0"

apply_radcal (bool, optional) - If set to True, will apply the pre-flight radiometric calibration curve found in the HDF5 header file and set units to $\text{erg}/(\text{cm}^2\text{ssr})$. If set to False, will simply return the data in units of photon counts. Default is True.

The return value is an `EISCube` class instance which contains calibrated intensities (or photon counts), corrected wavelengths, and all of the associated metadata. `EISCube` objects are a subclass of `NDCube` (from the Sunpy-affiliated package of the same name) and, as such, have built-in slicing and coordinate conversion capabilities due to an accompanying World Coordinate System (WCS) object. For example, you can slice an `NDCube` object using either array indices or by inputting physical coordinates to the `.crop_by_coords()` method. Please see the `ndcube` documentation³ for more information about slicing and manipulating `NDCube` objects.

The `EISCube` subclass extends `ndcube` by including a few additional features. First, an extra `.wavelength` attribute has been added which contains a 3D array with the corrected wavelength values at all locations within the cube. This correction accounts for a systematic spectral shift caused by a tilt in orientation of the EIS slit relative to the CCD. Slicing an `EISCube` will also appropriately slice the wave-

³ <https://docs.sunpy.org/projects/ndcube/en/stable/index.html>

length array. Secondly, four methods are included to quickly perform common EIS image processing,

- The `.apply_radcal()` and `.remove_radcal()` methods can be used to convert the data and uncertainty values to and from intensity and photon count units using the pre-flight radiometric calibration curve provided in the HDF5 header file. Currently, neither method takes any arguments. Future versions of EISPAC will allow users to specify their own calibration curves.
- The `.sum_spectra()` method sums the data along the wavelength axis and returns a new, 2D `NDCube` with just the data (no uncertainty or wavelength information). It requires no arguments.
- The `.smooth_cube()` method applies a boxcar moving average to the data along one or more spatial axes. It requires a single argument, "width", that must be either a singular value or list of ints, floats, or `astropy.units.Quantity` instances specifying the number of pixels or angular distance to smooth over. If given a single value, only the y-axis will be smoothed. Floats and angular distances will be converted to the nearest whole pixel value. If a width value is even, `width + 1` will be used instead. `.smooth_cube()` also accepts any number of optional keyword arguments that will be passed to the `astropy.convolution.convolve()` function, which does the actual smoothing operation.

The calibrated intensity and uncertainty values are stored in numpy arrays in the `.data` and `.uncertainty` attributes. The order of the axes are (slit position, raster step, wavelength) which correspond to the physical axes of (Solar-Y, Solar-X, Wavelength). You can inspect the dimensions of an `NDCube` object like so,

```
>>> data_cube.dimensions
[512, 87, 24] pix
```

As you can see, our example data has dimensions of (512, 87, 24). That is, 512 pixels along the slit (in the Solar-Y direction), 87 raster steps in the X direction, and 24 pixels in the dispersion direction.

All metadata and information from the HDF5 header file are packed into a single dictionary stored in the `.meta` attribute of the `EISCube`. The structure of the `.meta` dictionary mirrors the internal structure of the HDF5 file, with a few extra keys added for convenience. You can explore the contents with the usual Python commands,

```
>>> data_cube.meta.keys()
dict_keys(['filename_data', 'filename_head', 'wininfo', 'iwin', 'iwin_str',
          'index', 'pointing', 'wave', 'radcal', 'slit_width',
          'slit_width_units', 'ccd_offset', 'wave_corr', 'wave_corr_t',
```



```
'wave_corr_tilt', 'date_obs', 'date_obs_format', 'duration',
'duration_units', 'aspect_ratio', 'notes'])
>>> data_cube.meta['pointing']['x_scale']
2.9952
>>> data_cube.meta['radcal']
array([8.06751, 8.060929, 8.054517, 8.048271, 8.042198, 8.036295,
      8.030562, 8.024157, 8.017491, 8.010971, 8.0046015, 7.998385,
      7.9923196, 7.9864078, 7.980654, 7.975055, 7.969617, 7.9643393,
      7.959224, 7.9542727, 7.949487, 7.9448686, 7.9404206, 7.9361415],
      dtype=float32)
```

Here `x_scale` is the number of arcsec per step in the raster. Most EIS rasters take more than 1 arcsec per step, which degrades the spatial resolution but increases the cadence. The variable `radcal` is the pre-flight calibration curve for this data window. It includes all of the factors for converting counts directly to $\text{erg cm}^{-2} \text{s}^{-1} \text{sr}^{-1}$.

We can make a quick image of the EIS data by making use of the `.plot()` method provided in all `NDCube` objects (note, it usually helps to sum along the dispersion direction first).

```
>>> data_cube.sum_spectra().plot(aspect=data_cube.meta['aspect_ratio'])
```

The `.plot()` method can also be used to display the spectrum from a single pixel, as shown below. For illustration, we also convert the data back in units of photon counts (this is the same as dividing the calibrated data by the `.meta['radcal']` array).

```
>>> ix = 48
>>> iy = 326
>>> spec = data_cube[iy,ix,:].remove_radcal()
>>> spec_plot = spec.plot()
>>> spec_plot.set_title(f'ix = {ix}, iy = {iy}, units = counts')
```

To perform more advanced plotting, such as logarithmically scaling the intensities, you will need to extract the data from the `EISCube` and create the figure yourself using any of the various Python plotting libraries. For example,

```
import numpy as np
import matplotlib.pyplot as plt

raster_sum = np.sum(data_cube.data, axis=2) # or
              data_cube.sum_spectra().data
scaled_img = np.log10(raster_sum)

plt.figure()
plt.imshow(scaled_img, origin='lower',
           extent=data_cube.meta['extent_arcsec'], cmap='gray')
plt.title(data_cube.meta['date_obs'][-1])
plt.xlabel('Solar-X [arcsec]')
plt.ylabel('Solar-Y [arcsec]')
plt.show()
```




Figure 1: An example Fe XII 195.119 Å line profile from the raster.




Figure 2: An example image formed by summing the data for the Fe XII spectral window in the dispersion direction. In a subsequent chapter we'll discuss fitting the spectra.

Plotting tip: setting both "aspect" (`y_scale/x_scale`) and "extent" (data range as [left, right, bottom, top]) in `plt.imshow()` can sometimes give unexpected results. You may need to experiment with the combination of keywords needed to get the plot you expect.

We usually don't care about the numbering of the data windows. It's more natural to want to read the data corresponding to a particular wavelength. The `eispac.read_wininfo()` function can be used help identify the spectral contents of each data window. The function takes an input header file and returns a record array containing the window numbers, min and max wavelengths and primary spectral line for each data window. Note: for your convenience, a copy of the `wininfo` array is also stored in the `EISCube.meta` dictionary.

```
>>> import eispac
>>> header_filename = 'eis_20190404_131513.head.h5'
>>> wininfo = eispac.read_wininfo(header_filename, 195.12)
>>> wininfo.dtype.names
('iwin', 'line_id', 'wvl_min', 'wvl_max', 'nl', 'xs')
>>> wininfo[0:4]
rec.array([(0, 'Fe XI 180.400', 180.03426, 180.72559, 32, 661),
          (1, 'Ca XV 182.100', 181.75139, 182.44266, 32, 738),
          (2, 'Fe X 184.720', 183.82512, 185.5865 , 80, 831),
          (3, 'Fe XII 186.750', 186.3891 , 187.0802 , 32, 946)],
          dtype=[('iwin', '<i4'), ('line_id', '<U64'), ('wvl_min', '<f4'),
                  ('wvl_max', '<f4'), ('nl', '<i4'), ('xs', '<i4')])
```

We can then use a `numpy.where()` call on the `wininfo` array to map wavelength to window number. Users familiar with IDL may be interested to note that `numpy` record arrays can be accessed similarly to an IDL array of structures (e.g. instead of `wininfo['wvl_min']` below, you could also use `wininfo.wvl_min`).

```
>>> import numpy as np
>>> wvl = 195.119
>>> p = (wininfo['wvl_max'] - wvl)*(wvl - wininfo['wvl_min'])
>>> iwin = np.where(p >= 0)[0]
>>> iwin
array([7], dtype=int64)
```

If the result is an empty array, the wavelength is not in the data.

Fitting the Data

Fitting of the spectra involves selecting a spectral line of interest (e.g. Fe XII 195.12 Å) from one of the spectral windows of in the data, choosing a function (or combination of functions) to fit, and determining an initial guess for each parameter. The next ingredient for a fit is the selection of an optimization method. By default, EISPAC uses a Python implementation of the well-known IDL method **MPFIT**, which solves the non-linear least squares problem using the Levenberg- Marquardt algorithm. The Python module, `mpfit.py`, can be found on GitHub⁴ and is included in EISPAC. Future versions of the code will include full support for other fitting packages such as the newer `astropy.modeling` framework⁵.

Template files

In order to make fitting quick and easy, we've created a set of fit templates for different spectral lines. An `h5dump`⁶ on one of the template files shows that it contains a `/template` group for the initial guess on the fit parameters and a `/parinfo` group containing constraints on the parameters for use by `mpfit.py`.

```
h5dump -n fe_12_195_119.2c.template.h5
HDF5 "fe_12_195_119.2c.template.h5" {
FILE_CONTENTS {
  group /
  group /parinfo
  dataset /parinfo/fixed
  dataset /parinfo/limited
  dataset /parinfo/limits
  dataset /parinfo/tied
  dataset /parinfo/value
  group /template
  dataset /template/component
  dataset /template/data_e
  dataset /template/data_x
  dataset /template/data_y
  dataset /template/fit
  dataset /template/fit_back
  dataset /template/fit_gauss
  dataset /template/line_ids
  dataset /template/n_gauss
  dataset /template/n_poly
}
```

⁴ <https://github.com/segasai/astrolibpy/>

⁵ We have chosen to use MPFIT due to its flexibility and long (20+ year) legacy in solar and heliophysics. While newer fitting methods provide powerful tools for data exploration, they are still young and incur significant additional computation time (and are also more complicated to parallelize).

⁶ `h5dump` is a command line tool used to inspect the contents of an HDF5 file. It is included the Anaconda Python distribution platform, but can also be installed on its own.

```
dataset /template/order
dataset /template/wmax
dataset /template/wmin
}
```

The templates files are named according to the following pattern: {primary spectral line}.{number of Gaussians}c.template.h5. The function `eispac.read_template()` can be used to read a template file and examine the contents.

```
>>> import eispac
>>> tmplt_filename = 'fe_12_195_119.2c.template.h5'
>>> tmplt = eispac.read_template(tmplt_filename)
```

This produces the output below, which automatically calls the method `.print_parinfo()` to view the initial parameter values and constraints in a nice format (you can turn off the extra output by setting `quiet=True`).

```
Template file,
... fe_12_195_119.2c.template.h5
--- FIT TEMPLATE PARAMETER CONSTRAINTS ---
*      Value  Fixed  Limited  Limits      Tied
p[0]  57514.6647  0      1  0      0.0000  0.0000
p[1]   195.1179  0      1  1      195.0778 195.1581
p[2]    0.0289  0      1  1      0.0191  0.0510
p[3]  8013.4013  0      1  0      0.0000  0.0000
p[4]   195.1779  0      1  1      195.1378 195.2181  p[1]+0.06
p[5]    0.0289  0      1  1      0.0191  0.0510  p[2]
p[6]   664.3349  0      0  0      0.0000  0.0000
```

The structure of `parinfo` is specific to MPFIT and should be familiar to anyone who has used the original IDL version; please see the Appendix A for more details⁷. The templates provided with EIS-PAC consist of one or more Gaussian functions (with parameters in the order of peak, centroid, & width) followed by one or more background polynomial terms (usually just a single, constant value). The values `.template['n_gauss']` and `.template['n_poly']` indicate, respectively, the number of Gaussian functions and background polynomial terms in a given template.

Fitting spectra

Once you've read in a template file, you can use the central wavelength to find the desired spectral window in the data using `eispac.read_cube()`.

```
>>> data_filename = 'eis_20190404_131513.data.h5'
>>> data_cube = eispac.read_cube(data_filename, tmplt.central_wave)
```

As mentioned in the previous chapter, `read_cube()` automatically applies all of the pointing and wavelength corrections, bad data masking, and error estimations needed for scientific analysis. By default, the

⁷The `EISFitTemplate` object returned by `eispac.read_template` also generates a `.funcinfo` list. This list will help with the implementation of other fitting methods in the future, but is currently not used by the code and can, thereby, be safely ignored.

Reminder: The command line script `eis_fit_files` can be used to quickly fit a directory of files using one or more templates in another directory.

code also converts the data from photon counts to intensity units of $\text{erg cm}^{-2} \text{s}^{-1} \text{sr}^{-1}$ using the appropriate pre-flight calibration curve. This conversion can be disabled by setting the keyword `apply_radcal=False`, should you prefer to run your fits in count space.

On to the fitting! Now that you have a template and the data elements, you can perform a fit of the entire data cube by calling the top-level fitting routine, `eispac.fit_spectra()`⁸. The easiest way to use `fit_spectra()` is to just give it both an `EISCube` and `EISFitTemplate` object (or filepaths to the data and template HDF5 files). You may slice your `EISCube` however you wish before fitting and the code will loop over the data appropriately (this includes fitting a single spectra or slit observation). Additionally, `fit_spectra()` takes advantage of the multiprocessing package in the Python standard library to automatically parallelize the fitting process and minimize the run time. You may control the number of processing cores used for the fitting with `ncpu` keyword, or set it equal to "max" or None to use the maximum number of cores available⁹.

Here is a minimal example program that just loads and fits the data.

```
import matplotlib.pyplot as plt
import astropy.units as u
import eispac

if __name__ == '__main__':
    # input data and template files
    data_filepath = './eis_20190404_131513.data.h5'
    template_filepath = './fe_12_195_119.2c.template.h5'

    # read fit template
    tmplt = eispac.read_template(template_filepath)

    # Read spectral window into an EISCube
    data_cube = eispac.read_cube(data_filepath, tmplt.central_wave)

    # Fit the data, then save it to disk and test loading it back in
    fit_res = eispac.fit_spectra(data_cube, tmplt, ncpu='max')
    save_filepaths = eispac.save_fit(fit_res, save_dir='cwd')
    load_fit = eispac.read_fit(save_filepaths[0])
```

`fit_spectra()` outputs a `EISFitResult` object, which may be saved to and HDF5 file and read back in later using the `eispac.save_fit()` and `eispac.read_fit()` functions (as shown above). The full doc string for `fit_spectra()` can be found in Appendix B

The output fit parameters are stored in a dictionary of arrays.

```
>>> for key in fit_res.fit.keys():
...     print(key, fit_res.fit[key].dtype, fit_res.fit[key].shape)

line_ids <U14 (2,)
main_component int16 ()
n_gauss int16 ()
n_poly int16 ()
```

⁸Here's what's happening under the hood, `fit_spectra()` calls the helper function `scale_guess()` to scale the initial parameter values to the data, then `mpfit` is called to actually run the Levenberg-Marquardt fitting on a custom function that computes the deviates between the input spectrum and a multi-gaussian fit

⁹**IMPORTANT:** due to the specifics of how the multiprocessing library works, any statements that call `fit_spectra()` using `ncpu > 1` MUST be wrapped in a `"if __name__ == '__main__':"` statement in the top-level script or program. If such a "name guard" statement is not detected, `fit_spectra()` will fall back to using a single process. Unfortunately, this means you can not directly use parallel fitting from an interactive Python shell, you must first write a program that you save and run.

```

status float64 (126, 41)
chi2 float64 (126, 41)
wavelength float64 (126, 41, 24)
int float64 (126, 41, 2)
err_int float64 (126, 41, 2)
params float64 (126, 41, 7)
perror float64 (126, 41, 7)
component int32 (7,)
param_names <U32 (7,)

```

The `EISFitResult` object also has a few methods that make it easy to extract the fit parameters and compute the fit profiles. The use of these methods are demonstrated in the longer example program below, which also shows one way to select a data cutout. Please see Appendix B for more details about the `EISFitResult` methods.

```

import matplotlib.pyplot as plt
import astropy.units as u
import eispac

if __name__ == '__main__':
    # Read in the fit template and EIS observation
    data_filepath = './eis_20190404_131513.data.h5'
    template_filepath = './fe_12_195_119.2c.template.h5'
    tmplt = eispac.read_template(template_filepath)
    data_cube = eispac.read_cube(data_filepath, tmplt.central_wave)

    # Select a cutout of the raster (note the order of array & plot indices!)
    cutout_extent = [48, 165, 254, 378] # units of [arcsec]
    w_coords = data_cube.axis_world_coords('em.wl')
    lower_left = (cutout_extent[2]*u.arcsec, cutout_extent[0]*u.arcsec,
                  w_coords[0])
    upper_right = (cutout_extent[3]*u.arcsec, cutout_extent[1]*u.arcsec,
                  w_coords[-1])
    raster_cutout = data_cube.crop_by_coords(lower_left,
                                              upper_corner=upper_right)

    # Fit the data and save it to disk
    fit_res = eispac.fit_spectra(raster_cutout, tmplt, ncpu='max')
    save_filepaths = eispac.save_fit(fit_res, save_dir='cwd')

    # Extract array of total data and fit intensities
    sum_data_inten = raster_cutout.sum_spectra().data
    fit_wave_cube, fit_inten_cube = fit_res.get_fit_profile(component=[0,1])
    sum_fit_inten = fit_inten_cube.sum(axis=2)

    # Extract example fit profiles at a higher spectral resolution
    ex_coords = [43, 28] # [Y,X] array coords in units of [pixels]
    fit_x, fit_y = fit_res.get_fit_profile(coords=ex_coords,
                                           num_wavelengths=100)
    c0_fit_x, c0_fit_y = fit_res.get_fit_profile(component=0,
                                                  coords=ex_coords, num_wavelengths=100)
    c1_fit_x, c1_fit_y = fit_res.get_fit_profile(component=1,
                                                  coords=ex_coords, num_wavelengths=100)
    c2_fit_x, c2_fit_y = fit_res.get_fit_profile(component=2,
                                                  coords=ex_coords, num_wavelengths=100)
    sub_data = raster_cutout.data[ex_coords[0], ex_coords[1], :]
    sub_wave = raster_cutout.wavelength[ex_coords[0], ex_coords[1], :]
    sub_err = raster_cutout.uncertainty.array[ex_coords[0], ex_coords[1], :]

```

```

# Make a multi-panel figure with the cutout and example
fig = plt.figure()
plot_grid = fig.add_gridspec(nrows=2, ncols=2, hspace=0.5, wspace=0.3)

data_img = fig.add_subplot(plot_grid[0,0])
data_img.imshow(sum_data_inten, origin='lower', extent=cutout_extent,
                cmap='gray')
data_img.set_title('Data Cutout')
data_img.set_xlabel('Solar-X [arcsec]')
data_img.set_ylabel('Solar-Y [arcsec]')

fit_img = fig.add_subplot(plot_grid[0,1])
fit_img.imshow(sum_fit_inten, origin='lower', extent=cutout_extent,
               cmap='gray')
fit_img.set_title('Total Fit Intensity')
fit_img.set_xlabel('Solar-X [arcsec]')
fit_img.set_ylabel('Solar-Y [arcsec]')

profile = fig.add_subplot(plot_grid[1,:])
profile.errorbar(sub_wave, sub_data, yerr=sub_err,
                 ls='', marker='o', color='k')
profile.plot(fit_x, fit_y, color='b', label='Combined profile')
profile.plot(c0_fit_x, c0_fit_y, color='r', label='Gaussian 1')
profile.plot(c1_fit_x, c1_fit_y, color='r', ls='--', label='Gaussian 2')
profile.plot(c2_fit_x, c2_fit_y, color='g', label='Background')
profile.set_title(f'Cutout indices iy = {ex_coords[0]}, '
                 +f' ix = {ex_coords[1]}')
profile.set_xlabel('Wavelength [Å]')
profile.set_ylabel('Intensity [erg / (cm2 s sr)]')
profile.legend(loc='upper left')
plt.show()

```

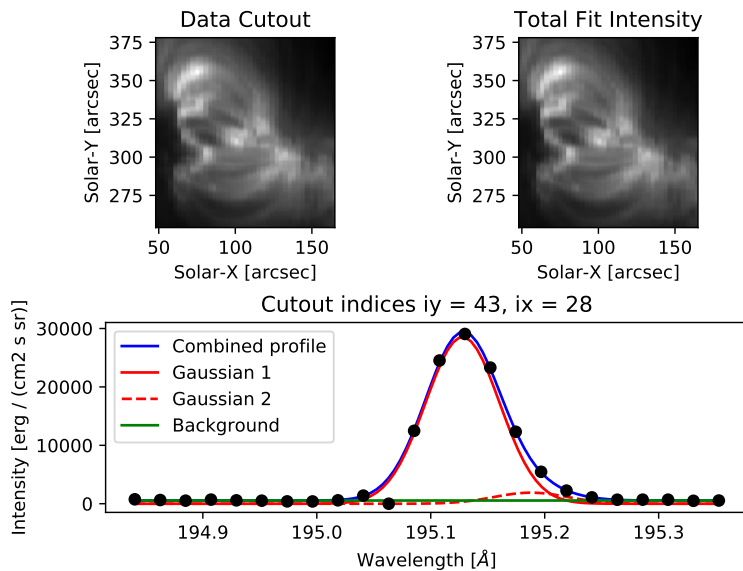


Figure 3: Example data cutout and profile fits. The top two panels show the raster formed by summing over the wavelength axis on the observed intensities (left) and fit intensities (right). The bottom panel shows an example fit for the Fe XII 195.119 Å line profile.

Level-1 HDF5 File Details

This chapter describes in more detail how the EIS level-1 HDF5 files were processed and saved. These HDF5 files can be downloaded from <https://eis.nrl.navy.mil/> or by using the search & download functionality of EISPAC (see Chapter 2).

Prepping the data in IDL

The level-0 fits files were prepped using the IDL routine `eis_prep` available via SolarSoft (Freeland & Handy, 1998) with the following options:

```
default = 1
save = 1
quiet = 1
retain = 1
photons = 1
refill = 0
```

There are 400,000+ EIS level-0 files at present, but on a multi-core machine using the IDL bridge all of the files can be prepped in under 24 hours. We have prepped all of the available EIS files and saved them to standard fits files in the usual way. Some important points:

units: As mentioned previously, the units for the output in these level-1 files is “photon events” or “counts.” This means that the statistical uncertainty can usually be estimated as \sqrt{N} . When EISPAC loads the HDF5 files, it also calculates an estimate for the read noise contribution. It should be noted, however, that the read noise becomes significant only at very low flux levels (1–2 counts).

retain: Note that the `retain` keyword preserves negative values. One of the jobs of `eis_prep` is to remove the pedestal from the CCD readout and any time-dependent dark current. Since the spectral windows are generally narrow, the estimate of the background can be too high and the subtracted intensities of the continuum can be negative. This will be dealt with during the fitting.

refill: The warm pixel problem complicates the fitting of EIS line profiles. As discussed in the EIS software note #13 (found in SSW or on the eiswiki), interpolating the values of missing pixels appears to best reproduce the original data. This option is left off during `eis_prep` so that the level-1 fits file preserves the information on the missing pixels. As discussed below, the interpolation (via the `refill` option) is done during the read and this data is ultimately written to the HDF5 file. A mask indicating which pixels have been interpolated will be added to the HDF5 files in a future revision.

Here is an IDL code snippet related to reading the data by looping over the spectral windows.

```
for iwin=0, nwin-1 do begin
  d = eis_getwindata(eis_level1_filename, iwin, /refill, /quiet)
  eis_level1_data[iwin] = ptr_new(d)
endfor
```

Writing the HDF5 files

Each processed level-1 fits file was bundled up with the associated calibration and metadata and saved as a pair of two HDF5 files:

eis_YYYYMMDD_HHMMSS.data.h5: Contains only the corrected photon counts within each spectral window. This is, by far, the larger of the two HDF5 files. However, they should not need to be updated or downloaded very often.

eis_YYYYMMDD_HHMMSS.head.h5: Contains the original fit file index, calibration curves for each spectral window (used to convert counts into intensity values), and the corrected pointing information.

Users will rarely, if ever, need to access the information inside the HDF5 files directly. EISPAC contains all of the functions needed to read the data and apply the calibration and pointing corrections. Nevertheless, the contents and data structure of the HDF5 files are summarized below.

.data.h5

level1 (group)

intensity_units (dataset) - String with the level1 data units. This will usually be "counts"

win## (dataset) - array of floating point values with the photon counts in a given spectral window (e.g. win00, win01, ... win24).

Note well, each set of EIS observations may have a different number of spectral windows, up to a maximum of 24 windows. Window numbers are numbered sequentially from 00; a given wavelength range may be assigned a different window number in each EIS study.

.head.h5

- *Details to be added soon*

Additionally, the contents of the HDF5 files can be displayed using h5dump command line tool, which is provided along with the Anaconda Python distribution platform or can be installed on its own. Example usage,

```
> h5dump -n eis_20190404_131513.data.h5
FILE_CONTENTS {
group    /
group    /level1
dataset  /level1/intensity_units
dataset  /level1/win00
dataset  /level1/win01
dataset  /level1/win02
dataset  /level1/win03
dataset  /level1/win04
dataset  /level1/win05
. . .
```

The actual data associated with each variable can be printed out using the -d option. For example,

```
> h5dump -d exposure_times/duration eis_20190404_131513.head.h5
HDF5 "eis_20190404_131513.head.h5" {
DATASET "exposure_times/duration" {
  DATATYPE H5T_IEEE_F32LE
  DATASPACE SIMPLE { ( 87 ) / ( 87 ) }
  DATA {
    (0): 40.0005, 40.0002, 40.0004, 40.0004, 39.9994, 40.0002, 39.9995, 40,
    (8): 40.0007, 39.9999, 40.0005, 40.0004, 39.9997, 40.0002, 39.9994,
    . . .
  }
}
```

Acknowledgments

This work was sponsored by NASA's Hinode project. Hinode is a Japanese mission developed and launched by ISAS/JAXA, with NAOJ as domestic partner and NASA and STFC (UK) as international partners.

Appendix A: mpfit documentation

parinfo keywords

Constraining Parameter Values with the PARINFO Keyword

The behavior of MPFIT can be modified with respect to each parameter to be fitted. A parameter value can be fixed; simple boundary constraints can be imposed; limitations on the parameter changes can be imposed; properties of the automatic derivative can be modified; and parameters can be tied to one another.

These properties are governed by the PARINFO structure, which is passed as a keyword parameter to MPFIT.

PARINFO should be a list of dictionaries, one list entry for each parameter. Each parameter is associated with one element of the array, in numerical order. The dictionary can have the following keys (none are required, keys are case insensitive):

'value' - the starting parameter value (but see the START_PARAMS parameter for more information).

'fixed' - a boolean value, whether the parameter is to be held fixed or not. Fixed parameters are not varied by MPFIT, but are passed on to MYFUNCT for evaluation.

'limited' - a two-element boolean array. If the first/second element is set, then the parameter is bounded on the lower/upper side. A parameter can be bounded on both sides. Both LIMITED and LIMITS must be given together.

'limits' - a two-element float array. Gives the parameter limits on the lower and upper sides, respectively. Zero, one or two of these values can be set, depending on the values of LIMITED. Both LIMITED and LIMITS must be given together.

'parname' - a string, giving the name of the parameter. The fitting code of MPFIT does not use this tag in any way. However, the default iterfunct will print the parameter name if available.

'step' - the step size to be used in calculating the numerical derivatives. If set to zero, then the step size is computed automatically.

Ignored when AUTODERIVATIVE=0.

'mpside' - the sidedness of the finite difference when computing numerical derivatives. This field can take four values:

0 - one-sided derivative computed automatically

1 - one-sided derivative $(f(x+h) - f(x))/h$

-1 - one-sided derivative $(f(x) - f(x-h))/h$

2 - two-sided derivative $(f(x+h) - f(x-h))/(2*h)$

Where "h" is the STEP parameter described above. The "automatic" one-sided derivative method will choose a direction for the finite difference which does not violate any constraints. The other methods do not perform this check. The two-sided method is in principle more precise, but requires twice as many function evaluations. Default: 0.

'mpmaxstep' - the maximum change to be made in the parameter value. During the fitting process, the parameter will never be changed by more than this value in one iteration. A value of 0 indicates no maximum. Default: 0.

'tied' - a string expression which "ties" the parameter to other free or fixed parameters. Any expression involving constants and the parameter array P are permitted. Example: if parameter 2 is always to be twice parameter 1 then use the following: `parinfo(2).tied = '2 * p(1)'`. Since they are totally constrained, tied parameters are considered to be fixed; no errors are computed for them. [NOTE: the PARNAME can't be used in expressions.]

'mpprint' - if set to 1, then the default iterfunc will print the parameter value. If set to 0, the parameter value will not be printed. This tag can be used to selectively print only a few parameter values out of many. Default: 1 (all parameters printed)

Appendix B: Useful EISPACE doc strings

fit_spectra()

```
def fit_spectra(inten, template, parinfo=None, wave=None, errs=None,
                min_points=10, ncpu='max', skip_fitting=False):
    """Fit one or more EIS line spectra using mpfit (with multiprocessing).

    Parameters
    -----
    inten : EISCube object, array_like, or filepath
        One or more intensity profiles to be fit. The code will loop over
        the data according to its dimensionality. 3D data is assumed to be a
        full EIS raster (or a sub region), 2D data is assumed to be a single
        EIS slit, and 1D data is assumed to be a single profile.
    template : EISFitTemplate object, dict, or filepath
        Either an EISFitTemplate, a 'template' dictionary, or the path to a
        template file.
    parinfo : list, optional
        List of dictionaries with fit parameters formatted for use with mpfit.
        Will supercede any parinfo lists loaded from an EISFitTemplate.
        Required if the 'template' parameter is given as a dictionary.
    wave : array_like, optional
        Associated wavelength values for the spectra. Required if 'inten' is
        given as an array and ignored otherwise.
    errs : array_like, optional
        Intensity error values for the spectra. Required if 'inten' is given
        as an array and ignored otherwise.
    min_points : int, optional
        Minimum number of good quality data points (i.e. non-zero values &
        errs) to be used in each fit. Spectra with fewer data points will be
        skipped. Default is 10.
    ncpu : int, optional
        Number of cpu processes to parallelize over. Must be less than or
        equal to the total number of cores the system has. If set to 'max' or
        None, the code will use the maximum number of cores available.
        Default is 'max'.
        Important: due to the specifics of how the multiprocessing library
        works, any statements that call fit_spectra() using ncpu > 1 MUST be
        wrapped in a "if __name__ == __main__:" statement in the top-level
        program. If such a "name guard" statement is not detected, this
        function will fall back to using a single process.
    skip_fitting : bool, optional
        If set to True, will skip the fitting altogether and just return an
        empty EISFitResult instance. Used mainly for testing.

    Returns
    -----
    fit_res : EISFitResult class instance
```

```

    An EISFitResult object containing the output fit parameters.
    """

```

EISFitResult methods

```

def get_params(self, component=None, param_name=None, coords=None,
               casefold=False):
    """Extract parameters values by component number, name, or pixel coords

    Parameters
    -----
    component : int or list, optional
        Integer number (or list of ints) of the functional component(s).
        If set to None, will return the total combined fit profile.
        Default is None.
    param_name : str, optional
        String name of the requested parameter. If set to None, will not
        filter based on parameter name. Default is None
    coords : list or tuple, optional
        (Y, X) coordinates of the requested datapoint. If set to None, will
        instead return the parameters at all locations. Default is None
    casefold : bool, optional
        If set to True, will ignore case when extracting parameters by
        name. Default is False.

    Returns
    -----
    param_vals : numpy array
        Parameter values
    param_errs : numpy array
        Estimated parameter errors
    """

```

```

def get_fit_profile(self, component=None, coords=None,
                   num_wavelengths=None):
    """Calculate the fit intensity profile (total or component) at a
    location.

    Parameters
    -----
    component : int or list, optional
        Integer number (or list of ints) of the functional component(s).
        If set to None, will return the total combined fit profile.
        Default is None.
    coords : list or tuple, optional
        (Y, X) coordinates of the requested datapoint. If set to None, will
        instead return the parameters at all locations. Default is None
    num_wavelengths : int, optional
        Number of wavelength values to compute the fit intensity at. These
        values will be equally spaced and span the entire fit window. If set
        to None, will use the observed wavelength values. Default is None.

    Returns
    -----
    fit_wave : numpy array
        Wavelength values

```

```
fit_inten : numpy array  
    Fit intensity values  
"""
```


Bibliography

Culhane, J. L., Harra, L. K., James, A. M., et al. 2007, *Sol. Phys.*, 243, 19

Freeland, S. L., & Handy, B. N. 1998, *Sol. Phys.*, 182, 497

Markwardt, C. B. 2009, in *Astronomical Society of the Pacific Conference Series*, Vol. 411, *Astronomical Data Analysis Software and Systems XVIII*, ed. D. A. Bohlender, D. Durand, & P. Dowler, 251